

Final Exam

CSE 202: Computer Organization and Architecture (Spring 2017)

May 15th, 2017

Name: _____ Lehigh Email ID: _____

*Instructions: Write down your name and Lehigh Email ID in the above spaces; you have 75 minutes; this is a closed-book, closed-notes exam; all calculators, PDAs, portable audio players and cell phones must be put away for the duration of the exam. **This exam is single-sided, no work/answers on the backs of any pages will be accepted.***

Write down your Lehigh Email ID at the top of each following page, do not put your name on those pages. Do that now. Seriously, don't screw that up. Your pages may not get graded if you do not do this.

| Exponent | 2^{Exponent} | Register | Usage |
|----------|-----------------------|--------------------|---------------------|
| 0 | 1 | rax | Return Value |
| 1 | 2 | rdi | Argument 0 |
| 2 | 4 | rsi | Argument 1 |
| 3 | 8 | rdx | Argument 2 |
| 4 | 16 | rcx | Argument 3 |
| 5 | 32 | r8 | Argument 4 |
| 6 | 64 | r9 | Argument 5 |
| 7 | 128 | rip | Instruction Pointer |
| 8 | 256 | rsp | Stack Pointer |
| 16 | 65536 | r10, r11 | Caller-saved Temps |
| 24 | 16777216 | rbx, r12, r13, r14 | Callee-saved Temps |
| 31 | 2147483648 | | |
| 32 | 4294967296 | | |

For named registers, prefix r is for 8 bytes, e is for 4 bytes.

For numbered registers, suffix d is for 4 bytes, no suffix for 8 bytes (r8-r14)

Floating Point: Sign Bit (s), Significant (M), Exponent (E) such that $-1^s * 2^E * M$ is the expressed value. Exponent E = unsigned exponent bits - bias. The bias is $2^{k-1} - 1$.

| Instruction | Meaning |
|-------------|--|
| mov | copies value into destination |
| lea | load effective address - does not dereference |
| call | goes to functions |
| cmp | subtracts and sets flags |
| jg | jump greater |
| ja | jump above (unsigned) |
| jmp | jump to an address |
| je | jump equal |
| jne | jump not equal |
| test | applies bitwise and to both arguments and sets flags |

Lehigh Email ID: _____

1. (10 points: 4/3/2/1/1, max 10): Shell Lab:

```
1 void sigchld_handler(int sig)
2 {
3     pid_t pid;
4     int status;
5
6     if((pid = waitpid(-1, &status, 0)) > 0)
7     {
8         if(WIFSTOPPED(status))
9         {
10             printf("Job [%d] (%d) stopped by signal %d\n",
11                 pid2jid(pid), pid, WSTOPSIG(status));
12             getjobpid(jobs, pid)->state = ST;
13             sigstp_handler(SIGTSTP);
14         }
15         else if(WIFSIGNALED(status))
16         {
17             printf("Job [%d] (%d) terminated by signal %d\n",
18                 pid2jid(pid), pid, WTERMSIG(status));
19             sigint_handler(SIGINT);
20         }
21         else
22         {
23             deletejob(jobs, pid);
24         }
25     }
26     if((errno != ECHILD) && (errno != 0))
27     {
28         unix_error("Waitpid error");
29     }
30 }
```

Consider the code above. There are at least five (5) distinct types of errors in this code. Indicate, by line number(s), what the five errors are and fix them (code fixes are best but an explanation is acceptable). Note that an error type can have multiple lines. Also, the same line can appear in more than one error type.

(a) Error 1 (4 points):

(b) Error 2 (3 points):

(c) Error 3 (2 points):

(d) Error 4 (1 point):

(e) Error 5 (1 point):

2. (10 points): Page Faults: When a page fault occurs, the kernel handles exception control flow in order to resolve the situation. What formally happens? Make sure to start from the task structures maintained by the kernel and consider how each possible outcome can be determined from the kernel. Do not forget to include how protection bits come into play here.

Specifically, assume you are the kernel and were just told that process X tried to access address Y and it was not valid in the page table.

Lehigh Email ID: _____

3. (10 points: 4/2/4): Malloc: Consider a situation where `realloc` is called on `ptr` in order to obtain a larger block. Assume the block that is physically after `ptr` is free but not of sufficient size for the reallocation.

If the block that is physically before `ptr` is sufficiently large (both on its own or in combination with the block physically after `ptr`), make an distinct argument for each of the following:

- (a) Reallocate into the block physically before `ptr` starting at the beginning of the block.
- (b) Reallocate into the block physically before `ptr` starting as far into the block as possible (leaving the beginning of that block free) and using the block physically afterwards, if applicable.
- (c) Just copy the block into a new, larger, location.

Lehigh Email ID: _____

4. (10 points) Malloc: Consider if we were to implement the free list(s) as Binary Search Tree(s) (BST). Describe all of the implementation and performance differences between a BST model and an explicit free list.

Lehigh Email ID: _____

5. (10 points): Process Creation: The system call *fork*, if implemented as described, is a rather expensive behavior in terms of number of cycles. In particular, *fork* is generally called with the intention of using *exec* afterwards. Explain how *demand paging* and *private copy-on-write (COW)* come together to allow us an efficient model for *fork/exec* semantics.

Lehigh Email ID: _____

6. (10 points: 7/3): Consider the code below. Assume that neither the compiler nor the CPU will reorder any of the code.

```
1 static jmp_buf env;
2 static int value = 4;
3
4 void foo(int *val)
5 {
6     value += (*val);
7     (*val)++;
8     printf("Value is now: %d\n", value);
9     longjmp(env, *val);
10 }
11
12 int main(int argc, char** argv)
13 {
14     int local = 24;
15     int jVal = (int)((setjmp(env) + local) / 3);
16     if(jVal < 10)
17     {
18         if(jVal & 0x1)
19         {
20             foo(&local);
21         }
22         else
23         {
24             foo(&value);
25         }
26     }
27     printf("Done, final values are: %d and %d\n", jVal, local);
28     return 0;
29 }
```

- (a) What is the output of this program?
- (b) Assume *local* is instead set to 1 when initialized. What is the output of this program?

Lehigh Email ID: _____

7. (10 points): I/O: Consider the code below:

```
1 int main(int argc , char **argv)
2 {
3     int fd1 , fd2 , fd3 , fd4;
4
5     fd1 = open(argv[1] , O_CREAT|O_TRUNC|O_RDWR, S_IRUSR|S_IWUSR);
6     write(fd1 , "apple" , 5);
7     fd2 = open(argv[1] , O_APPEND|O_WRONLY, 0);
8     write(fd2 , "banana" , 6);
9
10    fd3 = dup(fd1);
11    fd4 = dup(fd2);
12    write(fd4 , "carrot" , 6);
13
14    char buffer[10];
15    read(fd3 , &buffer[0] , 2);
16
17    write(fd3 , "donut" , 5);
18    write(fd4 , "eclair" , 6);
19
20    read(fd1 , &buffer[0] , 4);
21    write(fd3 , "fig" , 3);
22    return 0;
23 }
```

Assuming the code is executed with valid command line arguments answer the following. Assume that O APPEND will always append to the end (as if it calls lseek(file, 0, SEEK END) before each and every write) of the file. O CREAT—O TRUNC creates the file as a size 0 file.

- (a) What are the contents of the resulting file?
- (b) What is read by each read call in the code above? Specify each read by its line number.

Lehigh Email ID: _____

8. (10 points) Caches: Consider a 4-way set-associative cache that is initially empty with block sizes of 8 bytes, addresses that are 64 bits long and a total data capacity of 32KB. The following code performs very slowly and our boss is mad at us and doesn't trust us with a compiler any longer. Our only option to improve the performance is to redesign the cache itself. We have three options, we can a) double the block size, b) double the associativity (number of ways/columns) or c) double the number of lines (rows/sets). In each case, we are doubling the total data capacity of our cache to 64KB.

Which option will result in the best performance and why?

```
1 // assume array[0][0] is at address 0.
2 long long array[1 << 30][4];
3
4 for (size_t i=0; i<(1 << 30); i++)
5 {
6     for (size_t j=0; j<4; j++)
7     {
8         if ((i+1) < (1 << 30))
9         {
10             array[i][j] = array[i+1][j] + 1;
11         }
12     }
13 }
```

Lehigh Email ID: _____

9. (20 points): Virtual Memory

On some machine we are using, Virtual Addresses are 12 bits, Physical Addresses are 10 bits and the Page Size is 64 bytes. Given the following Page Table (incomplete but contains all that is relevant) and 2-way set-associative TLB, fill in the table specified below. Note that VPN is Virtual Page Number, PPN is Physical Page Number and V is valid. Assume all numbers are in hex.

| Page table | | | TLB | | | | | | |
|------------|-----|---|-------------|-------|-------|-------|-------|-------|-----|
| VPN | PPN | V | | Way A | | | Way B | | |
| 00 | 05 | 1 | Row | Tag | Valid | PPN | Tag | Valid | PPN |
| 01 | 01 | 0 | 0: | 03 | 0 | 0A | 02 | 1 | 0B |
| 02 | 07 | 0 | 1: | 03 | 0 | 0C | 02 | 0 | 0D |
| 03 | 0F | 1 | 2: | 04 | 1 | 03 | 00 | 1 | 07 |
| 04 | 0E | 1 | 3: | 03 | 1 | 0C | 02 | 1 | 0D |
| 05 | 04 | 1 | 4: | 01 | 0 | 06 | 03 | 1 | 06 |
| 06 | 0F | 0 | 5: | 03 | 1 | 0A | 02 | 0 | 0B |
| 07 | 07 | 1 | 6: | 05 | 1 | 09 | 03 | 1 | 07 |
| 08 | 06 | 0 | 7: | 01 | 1 | 06 | 03 | 0 | 06 |
| 09 | 03 | 0 | | | | | | | |
| 0A | 0C | 0 | | | | | | | |
| 0B | 04 | 1 | | 0x4FB | 0x61C | 0x333 | 0x52C | 0x002 | |
| 0C | 0C | 1 | VPN | | | | | | |
| 0D | 09 | 0 | PPN | | | | | | |
| 0E | 06 | 0 | TLB Tag | | | | | | |
| 0F | 0C | 1 | TLB Index | | | | | | |
| 10 | 0B | 0 | TLB Hit | | | | | | |
| 11 | 01 | 1 | Page Fault? | | | | | | |
| 12 | 02 | 0 | | | | | | | |
| 13 | 0D | 1 | | | | | | | |
| 14 | 03 | 1 | | | | | | | |
| 15 | 0A | 1 | | | | | | | |
| 16 | 0B | 1 | | | | | | | |
| 17 | 09 | 1 | | | | | | | |
| 18 | 0F | 1 | | | | | | | |

Did you remember to put your Lehigh Email ID at the top of each page? Did you?